


Aggregate error structure — all 35 models that learned the task

summed over 35 models × 90 annotations · macro AUC ≥ 0.9 · percentages are row-normalised



176 of 179 class confusions (98.3%) are thumbout ↔ closedhand. Only 3 involve openhand.

mean per-class AUC: thumbout 0.9497 openhands 0.9979 closedhand 0.9550

The 35 models are not independent trials — they are scored on the same 90 annotations, so an intrinsically ambiguous hand is counted repeatedly. The matrix shows the shape of the error reliably; read its magnitude as "how often across configurations", not as 3,150 independent observations.